

Genetic Algorithms

SEM6120 - Assignment 1

ALEXANDER D BROWN (ADB9)

Contents

1	Introduction	3
1.1	Evolutionary Algorithms	3
1.2	Genetic Algorithms	3
1.3	Chromosome Representation	4
1.4	Fitness Function	4
1.5	Selection	4
1.6	Crossover	5
1.7	Mutation	5
1.8	Termination	5
2	Previous Research into Selection Schemes	7
2.1	A Comparative Analysis of Selection Schemes used in Genetic Algorithms	7
2.1.1	Introduction	7
2.1.2	Comparison Techniques	7
2.1.3	Results	8
2.2	Genetic Drift in Selection Schemes	8
2.2.1	Introduction	8
2.2.2	Genetic Drift	8
2.2.3	Results	9
2.3	Sexual Selection with Competitive/Co-operative Operators for Genetic Algorithms	9
2.3.1	Introduction	9
2.3.2	Gendered Selection	10
2.3.3	Results	10
3	Evaluation of Research	11
3.1	Evaluation of A Comparative Analysis of Selection Schemes Used in Genetic Algorithms	11
3.2	Evaluation of Genetic Drift in Selection Schemes	11
3.3	Evaluation of Sexual Selection with Competitive/Co-operative operators for Genetic Algorithms	12

1 Introduction

Genetic algorithms are a biologically-inspired approach to heuristic search that mimic natural selection. Unlike many other evolutionary strategies and evolutionary programming, they are not designed to solve a specific problem, but are designed to solve the problem of optimisation which is made difficult by substantial complexity and uncertainty[2].

The complexity of the task should make it such that discovering an optimum solution is a long, maybe even impossible, task. At the same time the uncertainty needs to be reduced so that the knowledge of *available* options can be increased.

The initial design for a genetic algorithm was a method for moving from one population of chromosomes to another using a form of natural selection. This algorithm also included methods for crossover, mutation and inversion. This idea of having a large interbreeding population was the distinguishing feature from any past attempts which had only considered the parent and one offspring, where the offspring was simply a mutation of the parent[3].

1.1 Evolutionary Algorithms

As their name suggests, an evolutionary algorithm applies elements from the biological theory of evolution to the problem of optimisation. These elements include:

- Reproduction
- Mutation
- Recombination
- Selection

Typically, a population of candidate solutions are generated to which a fitness function can be applied. The population is then subject to some form of evolution, and this process is repeated until a halting criteria is met.

Genetic algorithms are a type of evolutionary algorithm with a focus on the genetic evolution of solutions. Candidate solutions for genetic algorithms, known as *chromosomes* are encoded as a series of *genes*. These genes are a representation of the choices which need to be optimised for the solution and can be as simple as a single bit or as complex as a real number, depending on the problem.

There are many other forms of both evolutionary and genetic algorithms which this report will mention in later sections.

1.2 Genetic Algorithms

The basic principals of genetic algorithms are to represent candidate solutions as a population of chromosomes, from this population the fittest members can be picked out and used in the next generation and to create new member of the population through reproduction and/or mutation.

This cycle repeats with the aim to produce better performing individuals in each generation until the optimum solution is either reached or gotten close

enough to that any future improvement is unnecessary or unwanted due to other constraints such as processing time. The latter of these allows a genetic algorithm to come up with a “good” solution in a reasonable amount of time.

Reproduction and mutation are an important part of genetic programming, and too of evolutionary programming. Without these parts the algorithm would quickly reach a local optimum for the initial population and would not improve past this.

1.3 Chromosome Representation

One of the key parts parts in implementing a genetic algorithm is the representation of chromosomes. This is very dependent of the problem the genetic algorithm needs to optimise and can have a knock on affect on the efficiency and accuracy of the algorithm.

Sometimes a simple solution is enough to represent the problem, binary strings are a commonly suggested approach. However sometimes more complex representations are required, potentially any data structure can be used as a chromosome but lists and trees are the common choices as they are easy to perform crossover¹ and mutation on.

As a very simple example, to maximise y in: $y = f(x)$, one could represent the value of x as a binary string, an example of which is shown in figure 1.

0	0	1	0	1	0	1	0
---	---	---	---	---	---	---	---

Figure 1: Chromosome representation as a binary string

1.4 Fitness Function

To fulfil the step of natural selection; the process of choosing the “best” members of a population, there needs to be a way of evaluating each chromosome, such that they can be compared to one another.

The function for doing so is known as a fitness function, which typically returns either a single number or a list of numbers, depending on the problem. Each chromosome can then be ranked in order of fitness and the top members of a population can then be selected.

As the value returned from a fitness function, with be specific to the domain it has be used within, it is necessary to rescale the fitness value to ensure uniformity when genetic algorithms are applied over several different domains at the same time.

1.5 Selection

Selection simulates the “survival of the fittest” nature of biology. However, it is beneficial not to keep some lesser performing members of the population to avoid getting trapped in local optimum. Most selection algorithms will introduce an element of randomness into the selection to deal with this.

¹A term used instead of reproduction in genetic algorithms.

Simple methods for selection typically involve randomly selecting individuals, giving those individuals with higher fitness greater chance of being selected. While others take inspiration from biology by making members of the population compete or assign individuals genders to imitate sexual reproduction.

1.6 Crossover

Crossover is the algorithmic approach to replicating biological reproduction and is key to genetic algorithms as it will produce new, unseen, offspring based on selected members of the population.

In its simplest form crossover is the process of taking one chromosome and splicing it with another. This can be done in several ways, but the most common is to choose a random point, known as the *crossover point*. The offspring consists of all the genes from the first parent up until the crossover point, the remaining genes then come from the second parent. As shown in figure 2.

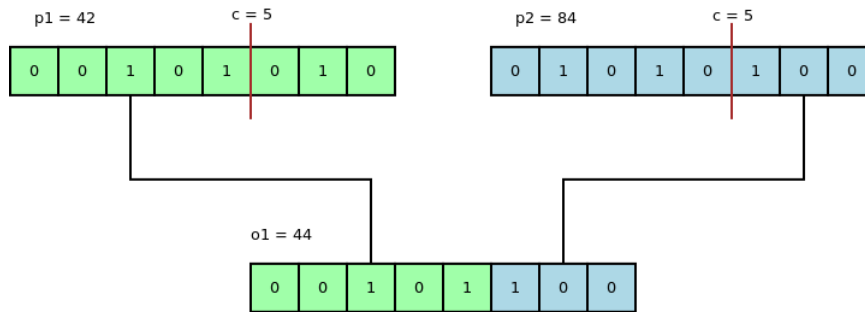


Figure 2: Crossover of parents $p1$ and $p2$ at crossover point c , where $c = 5$, resulting in offspring $o1$.

More complex methods revolve around the idea of having multiple crossover points and eventually it becomes more efficient to use a binary mask over the two parents, as shown in figure 3

This can lead to problems if the representation of chromosomes is complex; if, for example, each gene represents a node in the travelling salesman problem then crossover can produce routes which are invalid as nodes may be duplicated due to this process.

1.7 Mutation

Another method of getting variation into a population is mutation, the random change of genes within a chromosome. In genetic algorithms this is typically replicated by flipping a single bit in the chromosome.

As with crossover, mutation must be handled carefully with complex chromosome representations as it may lead to invalid chromosomes. The rate of mutation is usually kept low to avoid large amounts of variance.

1.8 Termination

The final problem with genetic algorithms is the problem of deciding when a result is good enough. Usually the implementation has some form of termination




Figure 3: Crossover of parents $p1$ and $p2$ using a binary mask, resulting in offspring $o1$

criteria, this may be a time limit or number of generations or may be something more complex like when the algorithm no longer improves in fitness greatly.

Termination criteria is very specific to the problem domain. Genetic algorithms are very good at giving an approximate answer and typically bad at finding the exact optimum solution.

2 Previous Research into Selection Schemes

This report focuses on reviewing previous research into selection schemes, focusing mainly on analysis of commonly used selection schemes through a range of methods, as well as viewing some more advanced techniques which are being used in recent years.

This section will briefly introduce each paper to give a general overview on each. The following section will then go on to critically evaluate each paper and start to draw conclusions on the research done.

2.1 A Comparative Analysis of Selection Schemes used in Genetic Algorithms

2.1.1 Introduction

Goldberg and Deb[1] were the first in the area to apply mathematical analysis to selection schemes. Through the application of four different analysis techniques they compared four methods of selection; proportional reproduction, ranking selection, tournament selection and the GENITOR algorithm[7].

From this they were able to perform side-by-side comparison of these four schemes and given useful information to those implementing genetic algorithms as to what selection schemes would be useful to which problem domains.

Though this is old research it sets the background for newer items and has sparked a lot of interest into the area. It also is still applicable today as complex selection schemes tend to be more domain specific, while older schemes, like tournament selection are still popular to this day.

2.1.2 Comparison Techniques

Selection schemes were analysed in terms of:

1. Growth ratio; the expected ratio for the members of the best class to the number of members of the population.
2. Takeover time; the approximate number of generations it takes to converge, this gives an idea of how long it would take before mutation, rather than crossover, becomes the primary method for exploring the search space.
3. Time complexity using the standard big O measure.

These values were chosen to allow side-by-side comparison of each selection scheme and help to provide useful information as to which scheme might be correct for a certain scenario.

The growth ratio is defined in three points of generation; 1) at generation t (ϕ_t), 2) when the proportion of individuals is insubstantial (the *early* ratio - ϕ_e) and 3) when the best individual occupies 50% of the population (the *late* ratio - ϕ_l).

These values are picked to give an idea of the performance in a general case for ϕ_t ; towards the beginning of convergence at ϕ_e , where growth rates are typically high and towards the end of convergence at ϕ_l where growth rates are expected to be low.

2.1.3 Results

Linear ranking and binary tournament (a tournament with a pool size of two individuals) were found to have similar growth rates. Tournaments selection could produce better growth rates simply through larger pool sizes, whereas ranking selection would need a non-linear function, making them more expensive to implement and run.

Takeover times all converged in around $O(\log n)$ generations. This is not at all surprising, given that the problem space only has a single optimum point. The authors note that this will not be the convergence of a real GA. The results of takeover time should be used to determine a rough estimate for when mutation, and not crossover, becomes the primary method of exploration.

Most interestingly the time complexities had a large range, from $O(n^2)$ to $O(n)$, with tournament selection being an $O(n)$. This is particularly interesting as tournament selection is a very easy algorithm to make parallel, this fits in with genetic algorithms nicely as they improve drastically from parallel implementation.

2.2 Genetic Drift in Selection Schemes

2.2.1 Introduction

Rogers and Prügel-Bennett[4] defined a method of analysing selection schemes using genetic drift, a term borrowed from biology, which describes the change in frequency of an allele (gene variation) through random sampling of the population.

Genetic drift is a phenomenon observed in genetic algorithms due to the nature of selection and, unlike analysis methods like convergence time, leads to a exact analytical solution. Though previous attempts to calculate genetic drift were often approximations or difficult to generalise to other cases, according to the authors.

2.2.2 Genetic Drift

In genetic algorithms, genetic drift is a measure of the change in fitness variation within a population and is worked out using probabilistic analysis of the selection scheme.

To analyse genetic drift, the variation of population fitness needs to be calculated. Standard variance (σ^2) is applied to the initial population of P members. Each member (α) has fitness F_α . Therefore the variance is defined as:

$$\sigma_F^2 = E[F^2] - (E[F])^2 \quad (1)$$

Calculating this out becomes:

$$\sigma_F^2 = \frac{1}{P} \sum_{\alpha=1}^P F_\alpha^2 - \left(\frac{1}{P} \sum_{\alpha=1}^P F_\alpha \right)^2 \quad (2)$$

Applying a selection scheme to this population and drawing a new population of P individuals means there will be n_α copies of a population member F_α . The variance of the new population is given as:

$$\sigma_F'^2 = \frac{1}{P} \sum_{\alpha=1}^P n_{\alpha} F_{\alpha}^2 - \left(\frac{1}{P} \sum_{\alpha=1}^P n_{\alpha} F_{\alpha} \right)^2 \quad (3)$$

To get an average case the authors averaged over all the possible ways of performing selection. To do this they looked at neutral selection, where the probability of selection is due to random, neutral, occurrences and not due to any form of fitness of the individual. This allows n_{α} to be independent of F_{α} . By working through this and simplifying further to use the fact that P is kept constant, they derive:

$$E[\sigma_F'^2] = \frac{P - E[n^2]}{P - 1} \sigma_F^2 \quad (4)$$

The change in population fitness variance for any selection scheme can be found by calculating $E[n^2]$, the expected square of the number of times any population member is selected.

By applying the variance in the number of times any individual is simply:

$$\sigma_n^2 = E[n^2] - (E[n])^2 \quad (5)$$

Substituting this into equation 4, the following equation is found, and is used as the basis for their results:

$$E[\sigma_F'^2] = \left(1 - \frac{\sigma_n^2}{P - 1} \right) \sigma_F^2 \quad (6)$$

The change in fitness variation due to selection is genetic drift. This is dependant only on the variance of the number of times any individual population member is selected (σ_n^2).

2.2.3 Results

To get comparable analysis a ratio r is defined as the change in variance after one generation.

$$r = \frac{E[\sigma_F'^2]}{\sigma_F^2} \quad (7)$$

$$= 1 - \frac{\sigma_n^2}{P - 1} \quad (8)$$

From this they concluded approximates of the genetic drift of several forms of selection scheme. Finding their results agreed with previous findings through empirical and numerical Markov chain comparisons[6].

2.3 Sexual Selection with Competitive/Co-operative Operators for Genetic Algorithms

2.3.1 Introduction

The idea assigning genders to chromosomes has been around for a long time and is typically done to restrict the crossover processes. The work of Sánchez-Velazco

and Bullinaria[5] expands on this to use the sexual characteristics of individuals as a part of selection, as well as being a part of crossover.

This process is designed to mimic more closely the process of natural evolution where there are real selection pressures, not only on individuals, but on mating pairs too.

2.3.2 Gendered Selection

As is normal for gendered genetic algorithms, individuals are assigned one of two genders: male or female. In the population, P , the set of females, X , and the set of males, Y , are such that $P = X \cup Y$ and $X \cap Y = \emptyset$.

Males are selected using a regular selection scheme, in this case roulette wheel selection was used. The selection function is defined as:

$$Sel(\phi(a)) \quad (9)$$

Where a is the set of all individuals with the fitness function ϕ . Choosing a male parent y_{sel} from $y \in Y$ is such that:

$$y_{sel} = Sel(f(y)) \quad (10)$$

The female parent is then selected from $x \in X$ according to a co-operative fitness function, which is based on: a) the competitive fitness f , b) the age of the female and c) the improvement in fitness Δf in the fitness of the offspring, compared to the selected male mate.

At the initial population, $\Delta f = 0$, all following generations will have a value for Δf computed for the female in the next selection process.

Choosing a female parent x_{sel} is such that:

$$x_{sel} = \frac{w_1 f(x) + w_2 \Delta f(y) + w_3 g(Age(x))}{w_1 + w_2 + w_3} \quad (11)$$

Where $\Delta f(y) = f(y_{son}) - f(y_{sel})$ and $g(Age(x))$ is a scaling function based on the age of the female parent.

2.3.3 Results

The authors used the Travelling Salesman Problem to test this method for selection and crossover. The results were averaged over 100 experimental runs with a population size of 200.

Comparing against a standard genetic algorithm in a problem with 20 cities, the gendered GA was found to have consistently better average fitness of both the whole population and the best individual. The latter of which reached the optimal fitness, while the standard GA only reached a fitness of $\frac{3}{4}$ the optimal.

Increasing the problem to have 50 cities, the gendered GA is found to have a much faster increase in fitness compared to a standard GA, again in terms of average fitness for the whole population and the best individual. More metrics were also applied to both the gendered and standard GAs to find out the resilience to noise and high mutation rates. Again the gendered GA outperforms the standard one, this is fairly intuitive as the selection strategy helps maintain a more "stable" population as the competition is only between individuals of the same gender.

3 Evaluation of Research

3.1 Evaluation of A Comparative Analysis of Selection Schemes Used in Genetic Algorithms

The results from Goldberg and Deb[1] should be taken with some scepticism as they are formulated from the pure mathematics of the selection techniques and not from running experiments and the authors even mention that it is only designed as a simple method to better understand the *expected* behaviour and suggest looking at better methods, citing the inherently noisy nature of genetic algorithms.

Its also notable that the results gained are not from a real scenario, rather a problem with a single local optimum and so results relating to convergence are not at all applicable to real genetic algorithms. No experimental results for real usage of the selection schemes is used to back up the results the authors gained.

Though a lot of the research done may not be applicable to “real” scenarios this work does help to provide some useful analysis to show what kinds of selection scheme might work well for a real problem; some problems might be able to trade off time complexity for a higher later growth rate and use the GENITOR algorithm, while other might require good time complexity to get a “good enough” solution quickly.

The main benefit of this work is that it has sparked off a lot more investigation into the performance of selection schemes in genetic algorithms and, although this work isn’t that substantial, it does provide a good starting point from which many other research have built upon.

Although there are a lot of flaw with this early research, it has provided a useful stepping stone for other research and, despite its age, does still provide a lot of useful information for those who plan to use selection schemes especially in the later sections which deal with the future of selection schemes.

3.2 Evaluation of Genetic Drift in Selection Schemes

The work of Rogers and Prügel-Bennett[4] also runs into this inherent problem of genetic algorithms being a form of subsymbolic learning, which is very difficult to apply probabilistic methods to, due to the randomness used within them.

As a counter-argument to this, it does give a good idea of the sorts of variation in fitness a selection scheme will provide and it does prove that genetic drift does affect the convergence of a genetic algorithm and can be controlled, at least to some degree, by the selection scheme.

Although the biological term has been modified to be applicable to genetic algorithms, it seems the authors have only taken into account the change in variance of fitness and not frequency of allele (gene variant) as the biological term is used to describe. Although this paper does have some nice ideas in ways in which selection schemes can be evaluated, it does not even start to broach the frequency of alleles in a genetic algorithm.

Of course, real genetic drift on a genetic algorithm may not be affected hugely by just the selection scheme, which is why the authors may have simplified the term.

Unlike the work of Goldberg and Deb, this paper does look at results from “real” genetic algorithms and compares the average over several thousand runs,

which does agree with their probabilistic findings. Again though, this should be taken with a little scepticism as there is little information on what sorts of problem domain(s) were used and it could be that the problems that the genetic algorithm was applied to was chosen to best match their results.

3.3 Evaluation of Sexual Selection with Competative/Co-operative operators for Genetic Algorithms

The first problem with the experimental results from the work of Sanchez-Velazco and Bullinaria[5] is the limited number of experimental runs performed. Even though the results are fairly conclusive in favour of the gendered genetic algorithm the random nature of genetic algorithms would leave some doubt as to whether 100 runs are enough to prove such a great enhancement.

Another noticeable flaw is the application of the selection scheme to only a single problem; the Travelling Salesman Problem is a very specialised problem for genetic algorithms and other forms of problems which require unordered chromosomes have been completely ignored.

This work does start to bridge the gap between genetic algorithms and other forms of evolutionary algorithms, particularly swarm-based algorithms like Particle Swarm Optimisation where the individuals have some knowledge of other individuals in the populations.

The results of this work are encouraging and the improvement, despite all the problems with the methods used to measure it, is impressive. For those applying GAs to the Travelling Salesman Problem this is definitely worth the time it would take to implement and should at least be at note for other problem domains.

The other potential issue with this technique is the addition of a second mutation rate, one for each gender. Rather than having to set a single mutation rate, there is an extra parameter which may add unneeded complexity to the algorithm.

All in all, it does seem like this work could have some really useful applications and, although it does seem like it would take a lot of work to find the right parameters for several parts of the algorithm, it definitely would be a useful addition to existing selection schemes and would not require much additional implementation to get it working on an existing system.

References

- [1] David E. Goldberg and Kalyanmoy Deb. A comparative analysis of selection schemes used in genetic algorithms. In *Foundations of Genetic Algorithms*, pages 69–93, 1991.
- [2] John H. Holland. *Adaptation in natural and artificial systems : an introductory analysis with applications to biology, control, and artificial intelligence*. MIT Press, April 1992.
- [3] Melanie Mitchell. An introduction to genetic algorithms, 1996.
- [4] A. Rogers and A. Prugel-Bennett. Genetic drift in genetic algorithm selection schemes. *Evolutionary Computation, IEEE Transactions on*, 3(4):298–303, November 1999.
- [5] Jose Sanchez-Velazco and John A. Bullinaria. Sexual selection with Competitive/co-Operative operators for genetic algorithms. In *In IASTED International Conference on Neural Networks and Computational Intelligence (NCI)*, 2003.
- [6] J. Schaffer, M. Mani, L. Eshelman, and K. Mathias. The effect of incest prevention on genetic drift. In W. Banzhaf and C. Reeves, editors, *Foundations of Genetic Algorithms*, volume 5, 1998.
- [7] Darrell Whitley. The GENITOR algorithm and selection pressure: Why Rank-Based allocation of reproductive trials is best. In *Proceedings of the Third International Conference on Genetic Algorithms*, pages 116–121, 1989.